

Homework 3

April 8, 2024

1 Homework 3

1.1 Section 1

```
[3]: # 1.1. Link: https://www.kaggle.com/datasets/uciml/default-of-credit-card-clients-dataset
      ↳ default-of-credit-card-clients-dataset
# 1.2. Description: This dataset contains information on default payments,
      ↳ demographic factors, credit data, history of payment, and bill statements of
      ↳ credit card clients in Taiwan from April 2005 to September 2005.
# 1.3. Fields/Attributes/Predictors:
# 1.4. Import Libraries
import numpy as np
import pandas as pd
import statsmodels.api as sm
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
from sklearn.model_selection import cross_val_score
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import KFold
from sklearn.metrics import mean_squared_error
from sklearn.model_selection import train_test_split
from sklearn.pipeline import make_pipeline
from sklearn.preprocessing import PolynomialFeatures
# 1.5. Load dataset into a Pandas dataframe
data = pd.read_csv('CreditCard.csv')
```

```
[4]: data.head(10)
```

```
[4]:   ID  LIMIT_BAL  SEX  EDUCATION  MARRIAGE  AGE  PAY_0  PAY_2  PAY_3  PAY_4  \
0    1    20000.0    2         2         1    24      2      2     -1     -1
1    2   120000.0    2         2         2    26     -1      2      0      0
2    3    90000.0    2         2         2    34      0      0      0      0
3    4    50000.0    2         2         1    37      0      0      0      0
4    5    50000.0    1         2         1    57     -1      0     -1      0
5    6    50000.0    1         1         2    37      0      0      0      0
6    7   500000.0    1         1         2    29      0      0      0      0
7    8   100000.0    2         2         2    23      0     -1     -1      0
```

8	9	140000.0	2	3	1	28	0	0	2	0
9	10	20000.0	1	3	2	35	-2	-2	-2	-2

	...	BILL_AMT4	BILL_AMT5	BILL_AMT6	PAY_AMT1	PAY_AMT2	PAY_AMT3	\
0	...	0.0	0.0	0.0	0.0	689.0	0.0	
1	...	3272.0	3455.0	3261.0	0.0	1000.0	1000.0	
2	...	14331.0	14948.0	15549.0	1518.0	1500.0	1000.0	
3	...	28314.0	28959.0	29547.0	2000.0	2019.0	1200.0	
4	...	20940.0	19146.0	19131.0	2000.0	36681.0	10000.0	
5	...	19394.0	19619.0	20024.0	2500.0	1815.0	657.0	
6	...	542653.0	483003.0	473944.0	55000.0	40000.0	38000.0	
7	...	221.0	-159.0	567.0	380.0	601.0	0.0	
8	...	12211.0	11793.0	3719.0	3329.0	0.0	432.0	
9	...	0.0	13007.0	13912.0	0.0	0.0	0.0	

	PAY_AMT4	PAY_AMT5	PAY_AMT6	default.payment.next.month
0	0.0	0.0	0.0	1
1	1000.0	0.0	2000.0	1
2	1000.0	1000.0	5000.0	0
3	1100.0	1069.0	1000.0	0
4	9000.0	689.0	679.0	0
5	1000.0	1000.0	800.0	0
6	20239.0	13750.0	13770.0	0
7	581.0	1687.0	1542.0	0
8	1000.0	1000.0	1000.0	0
9	13007.0	1122.0	0.0	0

[10 rows x 25 columns]

[5]: data.info

[5]: <bound method DataFrame.info of

	ID	LIMIT_BAL	SEX	EDUCATION	
MARRIAGE	AGE	PAY_0	PAY_2	PAY_3	\
0	1	20000.0	2	2	1 24 2 2 -1
1	2	120000.0	2	2	2 26 -1 2 0
2	3	90000.0	2	2	2 34 0 0 0
3	4	50000.0	2	2	1 37 0 0 0
4	5	50000.0	1	2	1 57 -1 0 -1
...
29995	29996	220000.0	1	3	1 39 0 0 0
29996	29997	150000.0	1	3	2 43 -1 -1 -1
29997	29998	30000.0	1	2	2 37 4 3 2
29998	29999	80000.0	1	3	1 41 1 -1 0
29999	30000	50000.0	1	2	1 46 0 0 0

	PAY_4	...	BILL_AMT4	BILL_AMT5	BILL_AMT6	PAY_AMT1	PAY_AMT2	\
0	-1	...	0.0	0.0	0.0	0.0	689.0	

1	0	...	3272.0	3455.0	3261.0	0.0	1000.0
2	0	...	14331.0	14948.0	15549.0	1518.0	1500.0
3	0	...	28314.0	28959.0	29547.0	2000.0	2019.0
4	0	...	20940.0	19146.0	19131.0	2000.0	36681.0
...
29995	0	...	88004.0	31237.0	15980.0	8500.0	20000.0
29996	-1	...	8979.0	5190.0	0.0	1837.0	3526.0
29997	-1	...	20878.0	20582.0	19357.0	0.0	0.0
29998	0	...	52774.0	11855.0	48944.0	85900.0	3409.0
29999	0	...	36535.0	32428.0	15313.0	2078.0	1800.0

	PAY_AMT3	PAY_AMT4	PAY_AMT5	PAY_AMT6	default.payment.next.month
0	0.0	0.0	0.0	0.0	1
1	1000.0	1000.0	0.0	2000.0	1
2	1000.0	1000.0	1000.0	5000.0	0
3	1200.0	1100.0	1069.0	1000.0	0
4	10000.0	9000.0	689.0	679.0	0
...
29995	5003.0	3047.0	5000.0	1000.0	0
29996	8998.0	129.0	0.0	0.0	0
29997	22000.0	4200.0	2000.0	3100.0	1
29998	1178.0	1926.0	52964.0	1804.0	1
29999	1430.0	1000.0	1000.0	1000.0	1

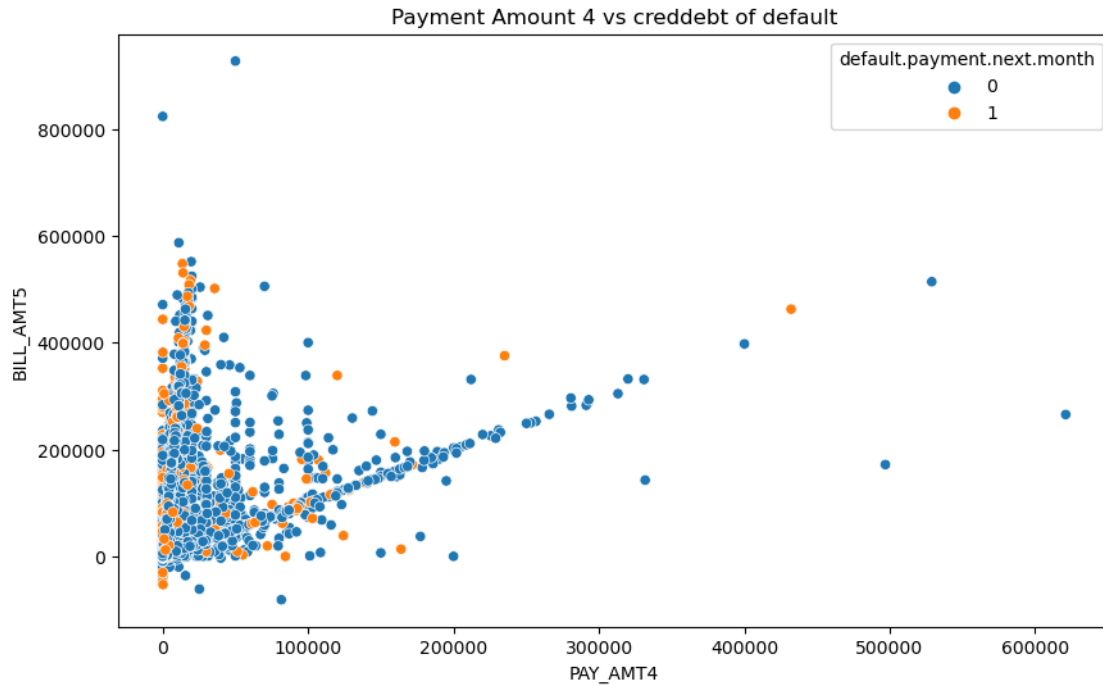
[30000 rows x 25 columns]>

1.2 Section 2

```
[7]: data = data.dropna(subset=['default.payment.next.month'])
X = data.loc[:, data.columns != "default.payment.next.month"]

X = sm.add_constant(X)
y = data['default.payment.next.month']

plt.figure(figsize=(10,6))
sns.scatterplot(x='PAY_AMT4', y='BILL_AMT5', hue='default.payment.next.month',
               data=data)
plt.title('Payment Amount 4 vs creddebt of default')
plt.show()
```



```
[8]: model = sm.Logit(y, X).fit()
      model.summary()
```

Optimization terminated successfully.
 Current function value: 0.464610
 Iterations 7

[8]:

Dep. Variable:	default.payment.next.month	No. Observations:	30000
Model:	Logit	Df Residuals:	29975
Method:	MLE	Df Model:	24
Date:	Mon, 08 Apr 2024	Pseudo R-squ.:	0.1208
Time:	19:32:48	Log-Likelihood:	-13938.
converged:	True	LL-Null:	-15853.
Covariance Type:	nonrobust	LLR p-value:	0.000

	coef	std err	z	P> z	[0.025	0.975]
const	-0.6675	0.121	-5.510	0.000	-0.905	-0.430
ID	-1.338e-06	1.75e-06	-0.765	0.444	-4.77e-06	2.09e-06
LIMIT_BAL	-7.615e-07	1.57e-07	-4.853	0.000	-1.07e-06	-4.54e-07
SEX	-0.1083	0.031	-3.530	0.000	-0.168	-0.048
EDUCATION	-0.1010	0.021	-4.815	0.000	-0.142	-0.060
MARRIAGE	-0.1548	0.032	-4.883	0.000	-0.217	-0.093
AGE	0.0074	0.002	4.170	0.000	0.004	0.011
PAY_0	0.5771	0.018	32.611	0.000	0.542	0.612
PAY_2	0.0832	0.020	4.119	0.000	0.044	0.123
PAY_3	0.0717	0.023	3.172	0.002	0.027	0.116
PAY_4	0.0248	0.025	0.990	0.322	-0.024	0.074
PAY_5	0.0334	0.027	1.240	0.215	-0.019	0.086
PAY_6	0.0080	0.022	0.361	0.718	-0.035	0.051
BILL_AMT1	-5.494e-06	1.14e-06	-4.836	0.000	-7.72e-06	-3.27e-06
BILL_AMT2	2.337e-06	1.51e-06	1.552	0.121	-6.14e-07	5.29e-06
BILL_AMT3	1.365e-06	1.32e-06	1.032	0.302	-1.23e-06	3.96e-06
BILL_AMT4	-8.861e-08	1.35e-06	-0.066	0.948	-2.74e-06	2.56e-06
BILL_AMT5	5.382e-07	1.52e-06	0.354	0.724	-2.45e-06	3.52e-06
BILL_AMT6	4.01e-07	1.2e-06	0.335	0.737	-1.94e-06	2.74e-06
PAY_AMT1	-1.363e-05	2.31e-06	-5.912	0.000	-1.82e-05	-9.11e-06
PAY_AMT2	-9.633e-06	2.09e-06	-4.599	0.000	-1.37e-05	-5.53e-06
PAY_AMT3	-2.723e-06	1.72e-06	-1.582	0.114	-6.1e-06	6.5e-07
PAY_AMT4	-3.967e-06	1.79e-06	-2.222	0.026	-7.47e-06	-4.68e-07
PAY_AMT5	-3.333e-06	1.78e-06	-1.874	0.061	-6.82e-06	1.52e-07
PAY_AMT6	-2.065e-06	1.3e-06	-1.593	0.111	-4.6e-06	4.75e-07

```
[9]: y_filtered = y[(data['AGE'] > 50) & (data['AGE'] < 60)]
y_filtered
data.iloc[[263]]
```

```
[9]:      ID  LIMIT_BAL  SEX  EDUCATION  MARRIAGE  AGE  PAY_0  PAY_2  PAY_3  \
263  264    230000.0    2           1           2   37     -2     -2     -2

      PAY_4  ...  BILL_AMT4  BILL_AMT5  BILL_AMT6  PAY_AMT1  PAY_AMT2  \
263     -2  ...         0.0       299.0       338.0   51315.0         0.0

      PAY_AMT3  PAY_AMT4  PAY_AMT5  PAY_AMT6  default.payment.next.month
263         0.0       299.0       338.0         0.0                      0

[1 rows x 25 columns]
```

```
[10]: test = X.iloc[[263]].copy()
test['AGE'] = 30
test['PAY_AMT4'] = 60
prediction = model.predict(test)
percentage_probability = prediction.iloc[0] * 100
```

```

binary_predictions = (prediction >= 0.5).astype(int)
class_prediction = binary_predictions.iloc[0]
print(f"Probability of default: {percentage_probability:.2f}%")
print(f"Class: {class_prediction}")
test

```

Probability of default: 3.11%

Class: 0

```

[10]:      const  ID  LIMIT_BAL  SEX  EDUCATION  MARRIAGE  AGE  PAY_0  PAY_2  \
263    1.0  264   230000.0    2           1           2   30    -2    -2

      PAY_3  ...  BILL_AMT3  BILL_AMT4  BILL_AMT5  BILL_AMT6  PAY_AMT1  \
263    -2  ...         0.0         0.0       299.0       338.0   51315.0

      PAY_AMT2  PAY_AMT3  PAY_AMT4  PAY_AMT5  PAY_AMT6
263         0.0         0.0        60      338.0         0.0

[1 rows x 25 columns]

```

1.3 Section 3

```

[12]: data['AGE'].unique()

```

```

[12]: array([24, 26, 34, 37, 57, 29, 23, 28, 35, 51, 41, 30, 49, 39, 40, 27, 47,
          33, 32, 54, 58, 22, 25, 31, 46, 42, 43, 45, 56, 44, 53, 38, 63, 36,
          52, 48, 55, 60, 50, 75, 61, 73, 59, 21, 67, 66, 62, 70, 72, 64, 65,
          71, 69, 68, 79, 74])

```

```

[13]: data['AGE'].value_counts()

```

```

[13]: AGE
29    1605
27    1477
28    1409
30    1395
26    1256
31    1217
25    1186
34    1162
32    1158
33    1146
24    1127
35    1113
36    1108
37    1041
39     954

```

38	944
23	931
40	870
41	824
42	794
44	700
43	670
45	617
46	570
22	560
47	501
48	466
49	452
50	411
51	340
53	325
52	304
54	247
55	209
56	178
58	122
57	122
59	83
60	67
21	67
61	56
62	44
63	31
64	31
66	25
65	24
67	16
69	15
70	10
68	5
73	4
72	3
75	3
71	3
79	1
74	1

Name: count, dtype: int64

```
[14]: X = data.loc[:, data.columns != "PAY_AMT4"] # Select all columns except "AGE"
X = sm.add_constant(X) # Add a constant column for the intercept
y = data['PAY_AMT4'] # Target variable
```

```

lda = LinearDiscriminantAnalysis()
lda.fit(X, y)

print(lda.explained_variance_ratio_)

X_r_lda = lda.transform(X)

target_names = ['1', '2', '3']
with plt.style.context('seaborn-talk'):
    fig, ax = plt.subplots(figsize=[10, 6])
    colors = ['blue', 'red', 'darkorange']

    # Plot the LDA results
    for color, i, target_name in zip(colors, [0, 1, 2], target_names):
        ax.scatter(X_r_lda[y == i, 0], X_r_lda[y == i, 1], alpha=0.8,
        ↪label=target_name, color=color)

    ax.set_title('LDA for Data')
    ax.set_xlabel('Discriminant Coordinate 1')
    ax.set_ylabel('Discriminant Coordinate 2')
    ax.legend(loc='best')


plt.show()

```

```

[0.1705034  0.10793117 0.08531915 0.07225079 0.06252642 0.05414688
 0.04740792 0.0426161  0.04091997 0.03668066 0.03298661 0.03198501
 0.02757561 0.02217313 0.02101852 0.02011324 0.01985403 0.01896578
 0.01745963 0.01638394 0.01520728 0.01399633 0.01314981 0.00882861]

```

1.4 Section 4

```
[16]: data = data.dropna(subset=['AGE'])
X = data.loc[:, data.columns != 'AGE']
X = sm.add_constant(X)
y = data['AGE']

lr = LinearRegression()

kf = KFold(n_splits=5, shuffle=True, random_state=42)
mse_scores = []

for train_index, test_index in kf.split(X):
    X_train, X_test = X.iloc[train_index], X.iloc[test_index]
    y_train, y_test = y.iloc[train_index], y.iloc[test_index]

    lr.fit(X_train, y_train)

    y_pred = lr.predict(X_test)

    mse = mean_squared_error(y_test, y_pred)
    mse_scores.append(mse)
    print(f"MSE for fold {len(mse_scores)}: {mse:.4f}")
```

```
print(f"Mean MSE from 5-fold CV using Linear Regression: {np.mean(mse_scores):.4f}")
print(f"Standard deviation of MSE: {np.std(mse_scores):.4f}")
```

```
MSE for fold 1: 66.0032
MSE for fold 2: 68.0571
MSE for fold 3: 67.4591
MSE for fold 4: 65.4976
MSE for fold 5: 66.6622
Mean MSE from 5-fold CV using Linear Regression: 66.7358
Standard deviation of MSE: 0.9321
```

1.5 Section 5

```
[18]: data = data.dropna(subset=['AGE'])
X = data.loc[:, data.columns != "AGE"]
X = sm.add_constant(X)
y = data['AGE']

def compute_cp(model, X, y):
    mse = np.mean((model.predict(X) - y) ** 2)
    p = len(model.params) - 1
    n = len(y)
    return mse + 2 * p * mse / (n - p - 1)

predictors = X.columns
cp_values, aic_values, bic_values, adjr2_values = [], [], [], []

for k in range(1, len(predictors) + 1):
    chosen_predictors = predictors[:k]
    X_subset = X[chosen_predictors]
    X_subset = sm.add_constant(X_subset)
    model = sm.OLS(y, X_subset).fit()
    cp_values.append(compute_cp(model, X_subset, y))
    aic_values.append(model.aic)
    bic_values.append(model.bic)
    adjr2_values.append(model.rsquared_adj)

fig, ax1 = plt.subplots(figsize=(12, 6))
ax2 = ax1.twinx()
ax3 = ax1.twinx()
ax4 = ax1.twinx()

ax3.spines['right'].set_position(('axes', 1.1))
ax4.spines['right'].set_position(('axes', 1.2))

ax1.plot(predictors, cp_values, 'g-', label="Cp", marker='o')
```

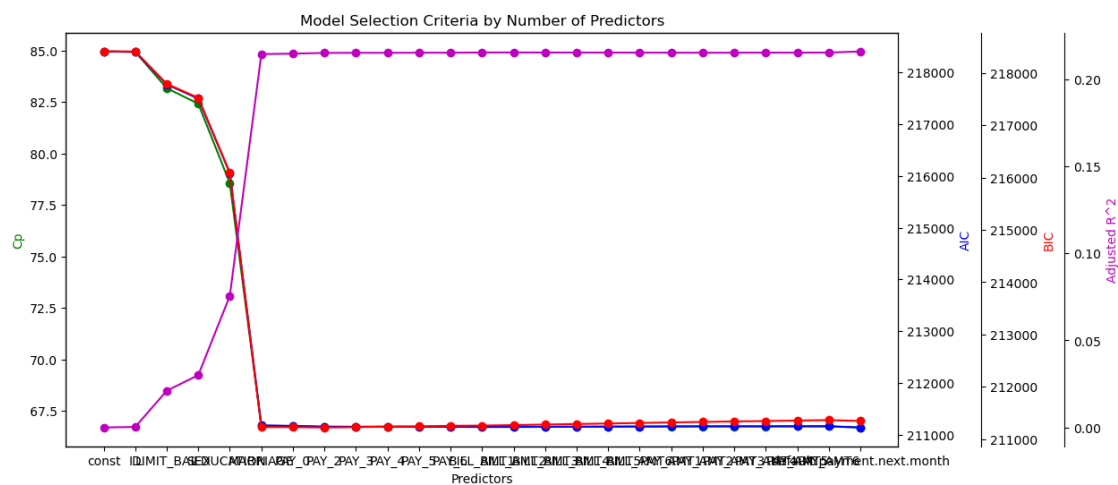
```

ax2.plot(predictors, aic_values, 'b-', label="AIC", marker='o')
ax3.plot(predictors, bic_values, 'r-', label="BIC", marker='o')
ax4.plot(predictors, adjr2_values, 'm-', label="Adj R^2", marker='o')

ax1.set_xlabel('Predictors')
ax1.set_ylabel('Cp', color='g')
ax2.set_ylabel('AIC', color='b')
ax3.set_ylabel('BIC', color='r')
ax4.set_ylabel('Adjusted R^2', color='m')

plt.title('Model Selection Criteria by Number of Predictors')
plt.show()

```



1.6 Section 6

```

[20]: data = data.dropna(subset=['AGE'])
X = data.loc[:, data.columns != "AGE"]
X = sm.add_constant(X)
y = data['AGE']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    random_state=42)

poly10_model = make_pipeline(PolynomialFeatures(3), LinearRegression())
poly10_model.fit(X_train, y_train)
y_pred_poly10 = poly10_model.predict(X_test)
mse_poly10 = mean_squared_error(y_test, y_pred_poly10)

poly4_model = make_pipeline(PolynomialFeatures(3), LinearRegression())
poly4_model.fit(X_train, y_train)

```

```
y_pred_poly4 = poly4_model.predict(X_test)
mse_poly4 = mean_squared_error(y_test, y_pred_poly4)

poly3_model = make_pipeline(PolynomialFeatures(3), LinearRegression())
poly3_model.fit(X_train, y_train)
y_pred_poly3 = poly3_model.predict(X_test)
mse_poly3 = mean_squared_error(y_test, y_pred_poly3)

print(f"Mean Squared Error for 10th-degree Polynomial: {mse_poly10:.2f}")
print(f"Mean Squared Error for 4th-degree Polynomial: {mse_poly4:.2f}")
print(f"Mean Squared Error for 3rd-degree Polynomial: {mse_poly3:.2f}")
```

Mean Squared Error for 10th-degree Polynomial: 542.78
Mean Squared Error for 4th-degree Polynomial: 542.78
Mean Squared Error for 3rd-degree Polynomial: 542.78